

SnowPro Core Practice Questions

Roushan Ara

August 19, 2025

Questions

Question 1

Snowflake provides a mechanism for its customers to override its natural clustering algorithms. This method is:

- A. Micro-partitions
- B. Clustering keys
- C. Key partitions
- D. Clustered partitions

See answer page 26

Question 2

Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)

- A. Custom
- B. Economy
- C. Optimized
- D. Standard

See answer page 26

Question 3

True or False: A single database can exist in more than one Snowflake account.

- A. True
- B. False

See answer page 26

Question 4

Which of the following roles is recommended to be used to create and manage users and roles?

- A. SYSADMIN
- B. SECURITYADMIN
- C. PUBLIC
- D. ACCOUNTADMIN

See answer page 26

Question 5

True or False: Bulk unloading of data from Snowflake supports the use of a SELECT statement.

- A. True
- B. False

See answer page 26

Question 6

Select the different types of Internal Stages: (Choose three.)

- A. Named Stage
- B. User Stage
- C. Table Stage
- D. Schema Stage

See answer page 26

Question 7

True or False: A customer using SnowSQL / native connectors will be unable to also use the Snowflake Web Interface (UI) unless access to the UI is explicitly granted by support.

- A. True
- B. False

See answer page 26

Question 8

Account-level storage usage can be monitored via:

- A. The Snowflake Web Interface (UI) in the Databases section
- B. The Snowflake Web Interface (UI) in the Account → Billing & Usage section
- C. The Information Schema → ACCOUNT_USAGE_HISTORY View
- D. The Account Usage Schema → ACCOUNT_USAGE_METRICS View

See answer page 26

Question 9

Credit Consumption by the Compute Layer (Virtual Warehouses) is based on: (Choose two.)

- A. Number of users
- B. Warehouse size
- C. Amount of data processed
- D. Number of clusters for the Warehouse

See answer page 26

Question 10

Which statement best describes ‘clustering‘?

- A. Clustering represents the way data is grouped together and stored within Snowflake’s micro-partitions
- B. The database administrator must define the clustering methodology for each Snowflake table
- C. The clustering key must be included on the COPY command when loading data into Snowflake
- D. Clustering can be disabled within a Snowflake account

See answer page 26

Question 11

True or False: The COPY command must specify a File Format in order to execute.

- A. True
- B. False

See answer page 26

Question 12

Which of the following commands sets the Virtual Warehouse for a session?

- A. COPY WAREHOUSE FROM <config file>
- B. SET WAREHOUSE = <warehouse name>
- C. USE WAREHOUSE <warehouse name>
- D. USE VIRTUAL_WAREHOUSE <warehouse name>

See answer page 27

Question 13

Which of the following objects can be cloned? (Choose four.)

- A. Tables
- B. Named File Formats
- C. Schemas
- D. Shares
- E. Databases
- F. Users

See answer page 27

Question 14

Which object allows you to limit the number of credits consumed within a Snowflake account?

- A. Account Usage Tracking
- B. Resource Monitor
- C. Warehouse Limit Parameter
- D. Credit Consumption Tracker

See answer page 27

Question 15

Snowflake is designed for which type of workloads? (Choose two.)

- A. OLAP (Analytics) workloads
- B. OLTP (Transactional) workloads
- C. Concurrent workloads
- D. On-premise workloads

See answer page 27

Question 16

What are the three layers that make up Snowflake's architecture? (Choose three.)

- A. Compute
- B. Tri-Secret Secure
- C. Storage
- D. Cloud Services

See answer page 27

Question 17

Why would a customer size a Virtual Warehouse from an X-Small to a Medium?

- A. To accommodate more queries
- B. To accommodate more users
- C. To accommodate fluctuations in workload
- D. To accommodate a more complex workload

See answer page 27

Question 18

True or False: Reader Accounts incur no additional Compute costs to the Data Provider since they are simply reading the shared data without making changes.

- A. True
- B. False

See answer page 27

Question 19

Which of the following connectors allow Multi-Factor Authentication (MFA) authorization when connecting? (Choose all that apply.)

- A. JDBC
- B. SnowSQL
- C. Snowflake Web Interface (UI)
- D. ODBC
- E. Python

See answer page 27

Question 20

True or False: Snowflake charges a premium for storing semi-structured data.

- A. True
- B. False

See answer page 27

Question 21

Which of the following statements describes a benefit of Snowflake's separation of compute and storage? (Choose all that apply.)

- A. Growth of storage and compute are tightly coupled together
- B. Storage expands without the requirement to add more compute
- C. Compute can be scaled up or down without the requirement to add more storage
- D. Multiple compute clusters can access stored data without contention

See answer page 27

Question 22

True or False: It is possible to unload structured data to semi-structured formats such as JSON and Parquet.

- A. True
- B. False

See answer page 27

Question 23

In which layer of its architecture does Snowflake store its metadata statistics?

- A. Storage Layer
- B. Compute Layer
- C. Database Layer
- D. Cloud Services Layer

See answer page 28

Question 24

True or False: Data in fail-safe can be deleted by a user or the Snowflake team before it expires.

- A. True
- B. False

See answer page 28

Question 25

True or False: Snowflake's data warehouse was built from the ground up for the cloud in lieu of using an existing database or a platform?

- A. True
- B. False

See answer page 28

Question 26

Which of the following statements are true of Virtual Warehouses? (Choose all that apply.)

- A. Customers can change the size of the Warehouse after creation
- B. A Warehouse can be resized while running
- C. A Warehouse can be configured to suspend after a period of inactivity
- D. A Warehouse can be configured to auto-resume when new queries are submitted

See answer page 28

Question 27

The PUT command: (Choose two.)

- A. Automatically creates a File Format object
- B. Automatically uses the last Stage created
- C. Automatically compresses files using Gzip
- D. Automatically encrypts files

See answer page 28

Question 28

Which type of table corresponds to a single Snowflake session?

- A. Temporary
- B. Transient
- C. Provisional
- D. Permanent

See answer page 28

Question 29

Which interfaces can be used to create and/or manage Virtual Warehouses?

- A. The Snowflake Web Interface (UI)
- B. SQL commands
- C. Data integration tools
- D. All of the above

See answer page 28

Question 30

When a Pipe is recreated using the CREATE OR REPLACE PIPE command:

- A. The Pipe load history is reset to empty
- B. The REFRESH parameter is set to TRUE
- C. Previously loaded files will be ignored
- D. All of the above

See answer page 28

Question 31

What is the minimum Snowflake edition that customers planning on storing protected information in Snowflake should consider for regulatory compliance?

- A. Standard
- B. Premier
- C. Enterprise
- D. Business Critical Edition

See answer page 28

Question 32

Select the three types of tables that exist within Snowflake. (Choose three.)

- A. Temporary
- B. Transient
- C. Provisional
- D. Permanent

See answer page 28

Question 33

True or False: Snowpipe via REST API can only reference External Stages as source.

- A. True
- B. False

See answer page 28

Question 34

True or False: A third-party tool that supports standard JDBC or ODBC but has no Snowflake-specific driver will be unable to connect to Snowflake.

- A. True
- B. False

See answer page 29

Question 35

True or False: It is possible to load data into Snowflake without creating a named File Format object.

- A. True
- B. False

See answer page 29

Question 36

True or False: A table in Snowflake can only be queried using the Virtual Warehouse that was used to load the data.

- A. True
- B. False

See answer page 29

Question 37

Which of the following statements are true of Snowflake data loading? (Choose three.)

- A. VARIANT null values are not the same as SQL NULL values
- B. It is recommended to do frequent, single row DMLs
- C. It is recommended to validate the data before loading into the Snowflake target table
- D. It is recommended to use staging tables to manage MERGE statements

See answer page 29

Question 38

Which statements are true of micro-partitions? (Choose two.)

- A. They are approximately 16MB in size
- B. They are stored compressed only if COMPRESS=TRUE on Table
- C. They are immutable
- D. They are only encrypted in the Enterprise edition and above

See answer page 29

Question 39

True or False: Query IDs are unique across all Snowflake deployments and can be used in communication with Snowflake Support to help troubleshoot issues.

- A. True
- B. False

See answer page 29

Question 40

A deterministic query is run at 8am, takes 5 minutes, and the results are cached. Which of the following statements are true? (Choose two.)

- A. The exact query will ALWAYS return the precomputed result set for the RESULT_CACHE_ACTIVE = time period
- B. The same exact query will return the precomputed results if the underlying data hasnt changed and the results were last accessed within previous 24 hour period
- C. The same exact query will return the precomputed results even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period
- D. The 24 hour timer on the precomputed results gets renewed every time the exact query is executed

See answer page 29

Question 41

Increasing the maximum number of clusters in a Multi-Cluster Warehouse is an example of:

- A. Scaling rhythmically
- B. Scaling max
- C. Scaling out
- D. Scaling up

See answer page 29

Question 42

Which statement best describes Snowflake tables?

- A. Snowflake tables are logical representations of underlying physical data
- B. Snowflake tables are the physical instantiation of data loaded into Snowflake
- C. Snowflake tables require that clustering keys be defined to perform optimally
- D. Snowflake tables are owned by a user

See answer page 29

Question 43

Which item in the Data Warehouse migration process does not apply in Snowflake?

- A. Migrate Users
- B. Migrate Schemas
- C. Migrate Indexes
- D. Build the Data Pipeline

See answer page 29

Question 44

Snowflake provides two mechanisms to reduce data storage costs for short-lived tables. These mechanisms are: (Choose two.)

- A. Temporary Tables
- B. Transient Tables
- C. Provisional Tables
- D. Permanent Tables

See answer page 30

Question 45

What is the maximum compressed row size in Snowflake?

- A. 8KB
- B. 16MB
- C. 50MB
- D. 4000GB

See answer page 30

Question 46

Which of the following are main sections of the top navigation of the Snowflake Web Interface (UI)? (Choose three.)

- A. Databases
- B. Tables
- C. Warehouses
- D. Worksheets

See answer page 30

Question 47

What is the recommended Snowflake data type to store semi-structured data like JSON?

- A. VARCHAR
- B. RAW
- C. LOB
- D. VARIANT

See answer page 30

Question 48

Which of the following statements are true of Snowflake releases: (Choose two.)

- A. They happen approximately weekly
- B. They roll up and release approximately monthly, but customers can request early release application
- C. During a release, new customer requests/queries/connections transparently move over to the newer version
- D. A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and customer is upgraded

See answer page 30

Question 49

Which of the following are common use cases for zero-copy cloning? (Choose three.)

- A. Quick provisioning of Dev and Test/QA environments
- B. Data backups
- C. Point in time snapshots
- D. Performance optimization

See answer page 30

Question 50

If a Small Warehouse is made up of 2 servers/cluster, how many servers/cluster make up a Medium Warehouse?

- A. 4
- B. 16
- C. 32
- D. 128

See answer page 30

Question 51

True or False: When a data share is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data consumers.

- A. True
- B. False

See answer page 30

Question 52

Which is true of Snowflake network policies? A Snowflake network policy: (Choose two.)

- A. Is available to all Snowflake Editions
- B. Is only available to customers with Business Critical Edition
- C. Restricts or enables access to specific IP addresses
- D. Is activated using an ALTER DATABASE command

See answer page 30

Question 53

True or False: Snowflake charges additional fees to Data Providers for each Share they create.

- A. True
- B. False

See answer page 30

Question 54

Query results are stored in the Result Cache for how long after they are last accessed, assuming no data changes have occurred?

- A. 1 Hour
- B. 3 Hours
- C. 12 Hours
- D. 24 Hours

See answer page 31

Question 55

A role is created and owns 2 tables. This role is then dropped. Who will now own the two tables?

- A. The tables are now orphaned
- B. The user that deleted the role
- C. SYSADMIN
- D. The assumed role that dropped the role

See answer page 31

Question 56

Which of the following connectors are available in the Downloads section of the Snowflake Web Interface (UI)? (Choose two.)

- A. SnowSQL
- B. ODBC
- C. R
- D. HIVE

See answer page 31

Question 57

Which of the following *DML* commands isn't supported by Snowflake?

- A. UPSERT
- B. MERGE
- C. UPDATE
- D. TRUNCATE TABLE

See answer page 31

Question 58

Which of the following statements is true of zero-copy cloning?

- A. Zero-copy clones increase storage costs as cloning the table requires storing its data twice
- B. All zero-copy clone objects inherit the privileges of their original objects
- C. Zero-copy cloning is licensed as an additional Snowflake feature
- D. At the instance/instant a clone is created, all micro-partitions in the original table and the clone are fully shared

See answer page 31

Question 59

True or False: When a user creates a role, they are initially assigned ownership of the role and they maintain ownership until it is transferred to another user.

- A. True
- B. False

See answer page 31

Question 60

The Query History in the Snowflake Web Interface (UI) is kept for approximately:

- A. 60 minutes
- B. 24 hours
- C. 14 days
- D. 30 days
- E. 1 year

See answer page 31

Question 61

To run a Multi-Cluster Warehouse in auto-scale mode, a user would:

- A. Configure the Maximum Clusters setting to Auto-Scale
- B. Set the Warehouse type to Auto
- C. Set the Minimum Clusters and Maximum Clusters settings to the same value
- D. Set the Minimum Clusters and Maximum Clusters settings to different values

See answer page 31

Question 62

Which of the following terms best describes Snowflake's database architecture?

- A. Columnar shared nothing
- B. Shared disk
- C. Multi-cluster, shared data
- D. Cloud-native shared memory

See answer page 31

Question 63

Which of the following are options when creating a Virtual Warehouse? (Choose two.)

- A. Auto-drop
- B. Auto-resize
- C. Auto-resume
- D. Auto-suspend

See answer page 31

Question 64

A Virtual Warehouse's auto-suspend and auto-resume settings apply to:

- A. The primary cluster in the Virtual Warehouse
- B. The entire Virtual Warehouse
- C. The database the Virtual Warehouse resides in
- D. The queries currently being run by the Virtual Warehouse

See answer page 32

Question 65

Fail-safe is unavailable on which table types? (Choose two.)

- A. Temporary
- B. Transient
- C. Provisional
- D. Permanent

See answer page 32

Question 66

Which of the following objects is not covered by Time Travel?

- A. Tables
- B. Schemas
- C. Databases
- D. Stages

See answer page 32

Question 67

True or False: Micro-partition metadata enables some operations to be completed without requiring Compute.

- A. True
- B. False

See answer page 32

Question 68

Which of the following commands are not blocking operations? (Choose two.)

- A. UPDATE
- B. INSERT
- C. MERGE
- D. COPY

See answer page 32

Question 69

Which of the following is true of Snowpipe via REST API? (Choose two.)

- A. You can only use it on Internal Stages
- B. All COPY INTO options are available during pipe creation
- C. Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands
- D. Snowpipe keeps track of which files it has loaded

See answer page 32

Question 70

Snowflake recommends, as a minimum, that all users with the following role(s) should be enrolled in Multi-Factor Authentication (MFA):

- A. SECURITYADMIN, ACCOUNTADMIN, PUBLIC, SYSADMIN
- B. SECURITYADMIN, ACCOUNTADMIN, SYSADMIN
- C. SECURITYADMIN, ACCOUNTADMIN
- D. ACCOUNTADMIN

See answer page 32

Question 71

When can a Virtual Warehouse start running queries?

- A. 12am–5am
- B. Only during administrator defined time slots
- C. When its provisioning is complete
- D. After replication

See answer page 32

Question 72

True or False: Users are able to see the result sets of queries executed by other users that share their same role.

- A. True
- B. False

See answer page 32

Question 73

True or False: The user must specify which cluster a query will run on in a multi-cluster Warehouse.

- A. True
- B. False

See answer page 32

Question 74

True or False: Pipes can be suspended and resumed.

- A. True
- B. False

See answer page 33

Question 75

Which of the following languages can be used to implement Snowflake User Defined Functions (UDFs)? (Choose two.)

- A. Java
- B. Javascript
- C. SQL
- D. Python

See answer page 33

Question 76

When should you consider disabling auto-suspend for a Virtual Warehouse? (Choose two.)

- A. When users will be using compute at different times throughout a 24/7 period
- B. When managing a steady workload
- C. When the compute must be available with no delay or lag time
- D. When you do not want to have to manually turn on the Warehouse each time a user needs it

See answer page 33

Question 77

Which of the following are valid approaches to loading data into a Snowflake table? (Choose all that apply.)

- A. Bulk copy from an External Stage
- B. Continuous load using Snowpipe REST API
- C. The Snowflake Web Interface (UI) data loading wizard
- D. Bulk copy from an Internal Stage

See answer page 33

Question 78

If auto-suspend is enabled for a Virtual Warehouse, the Warehouse is automatically suspended when:

- A. All Snowflake sessions using the Warehouse are terminated
- B. The last query using the Warehouse completes
- C. There are no users logged into Snowflake

See answer page 33

Question 79

True or False: Multi-Factor Authentication (MFA) in Snowflake is only supported in conjunction with Single Sign-On (SSO).

- A. True
- B. False

See answer page 33

Question 80

The number of queries that a Virtual Warehouse can concurrently process is determined by (Choose two.):

- A. The complexity of each query
- B. The `CONCURRENT_QUERY_LIMIT` parameter set on the Snowflake account
- C. The size of the data required for each query
- D. The tool that is executing the query

See answer page 33

Question 81

Which of the following statements are true of `VALIDATION_MODE` in Snowflake? (Choose two.)

- A. The `VALIDATION_MODE` option is used when creating an Internal Stage
- B. `VALIDATION_MODE=RETURN_ALL_ERRORS` is a parameter of the `COPY` command
- C. The `VALIDATION_MODE` option will validate data to be loaded by the `COPY` statement while completing the load and will return the rows that could not be loaded without error
- D. The `VALIDATION_MODE` option will validate data to be loaded by the `COPY` statement without completing the load and will return possible errors

See answer page 33

Question 82

What privileges are required to create a task?

- A. The `GLOBAL` privilege `CREATE TASK` is required to create a new task.
- B. Tasks are created at the Application level and can only be created by the Account Admin role.
- C. Many Snowflake DDLs are metadata operations only, and `CREATE TASK` DDL can be executed without virtual warehouse requirement or task specific grants.
- D. The role must have access to the target schema and the `CREATE TASK` privilege on the schema itself.

See answer page 33

Question 83

What are the three things customers want most from their enterprise data warehouse solution? (Choose three.)

- A. On-premise availability
- B. **Simplicity**
- C. Open source based
- D. Concurrency
- E. **Performance**

See answer page 34

Question 84

True or False: Some queries can be answered through the metadata cache and do not require an active Virtual Warehouse.

- A. True
- B. False

See answer page 34

Question 85

When scaling out by adding clusters to a multi-cluster warehouse, you are primarily scaling for improved:

- A. Concurrency
- B. Performance

See answer page 34

Question 86

What is the minimum Snowflake edition that provides data sharing?

- A. Standard
- B. Premier
- C. Enterprise
- D. Business Critical Edition

See answer page 34

Question 87

True or False: Each worksheet in the Snowflake Web Interface (UI) can be associated with different roles, databases, schemas, and Virtual Warehouses.

A. True

B. False

See answer page 34

Question 88

True or False: You can query the files in an External Stage direc

Answers

Question 1: B. Clustering keys

They allow customers to explicitly define clustering for performance optimization.

[Back to Question 1](#)

Question 2: B. Economy **and** D. Standard

These are the only valid Snowflake Virtual Warehouse scaling policies.

[Back to Question 2](#)

Question 3: B. False

A Snowflake database belongs to a single account and cannot exist in more than one account.

[Back to Question 3](#)

Question 4: B. SECURITYADMIN

This role is specifically responsible for creating and managing users and roles.

[Back to Question 4](#)

Question 5: A. True

Bulk unloading supports the use of a SELECT statement to filter and transform data before writing to files.

[Back to Question 5](#)

Question 6: A. Named Stage, B. User Stage, and C. Table Stage

These are the three types of internal stages in Snowflake.

[Back to Question 6](#)

Question 7: B. False

SnowSQL and native connectors can be used independently of the Web UI; UI access does not require explicit support enablement.

[Back to Question 7](#)

Question 8: D. The Account Usage Schema → ACCOUNT_USAGE_METRICS View

This view provides account-level storage usage details.

[Back to Question 8](#)

Question 9: B. Warehouse size **and** D. Number of clusters for the Warehouse

Credit consumption depends on warehouse size and the number of clusters in multi-cluster warehouses.

[Back to Question 9](#)

Question 10: A. Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions

Clustering is Snowflake's method of organizing micro-partition storage.

[Back to Question 10](#)

Question 11: A. True

The COPY command requires a file format (or defaults) to execute properly.

[Back to Question 11](#)

Question 12: C. USE WAREHOUSE <warehouse name>

This command sets the active virtual warehouse for the session.

[Back to Question 12](#)

Question 13: A. Tables, C. Schemas, D. Shares, E. Databases

These Snowflake objects can be cloned. Named file formats and users cannot.

[Back to Question 13](#)

Question 14: B. Resource Monitor

Resource monitors are used to limit and control credit consumption within a Snowflake account.

[Back to Question 14](#)

Question 15: A. OLAP (Analytics) workloads **and** C. Concurrent workloads

Snowflake is designed for analytics and concurrent workloads, not OLTP or on-premise.

[Back to Question 15](#)

Question 16: A. Compute, C. Storage, D. Cloud Services

These are the three layers of Snowflakes architecture.

[Back to Question 16](#)

Question 17: B. To accommodate more users

Larger virtual warehouses support more users running queries concurrently.

[Back to Question 17](#)

Question 18: B. False

Reader Accounts still generate compute costs for the provider since queries executed on shared data consume compute resources.

[Back to Question 18](#)

Question 19: C. Snowflake Web Interface (UI) and B. SnowSQL

MFA is supported through the web interface and SnowSQL client, but not directly via JDBC/ODBC/Python.

[Back to Question 19](#)

Question 20: B. False

Snowflake does not charge a premium for semi-structured data; it is stored in VARIANT columns without extra fees beyond normal storage.

[Back to Question 20](#)

Question 21: B. Storage expands without adding compute, C. Compute scales independently of storage, D. Multiple clusters can access data without contention

These are the key benefits of separating compute from storage in Snowflake.

[Back to Question 21](#)

Question 22: A. True

Snowflake supports unloading structured data into semi-structured formats such as JSON and Parquet.

[Back to Question 22](#)

Question 23: D. Cloud Services Layer

Snowflake stores and manages metadata statistics in the Cloud Services layer.

[Back to Question 23](#)

Question 24: B. False

Fail-safe is managed entirely by Snowflake and cannot be modified or deleted by users or Snowflake staff before expiration.

[Back to Question 24](#)

Question 25: A. True

Snowflake was designed from the ground up as a cloud-native data warehouse platform, not adapted from an on-premise system.

[Back to Question 25](#)

Question 26: A. Customers can change size after creation, C. Suspend after inactivity, D. Auto-resume on query submission

Warehouses cannot be resized dynamically while running (option B is false).

[Back to Question 26](#)

Question 27: C. Automatically compresses files using Gzip, D. Automatically encrypts files
The PUT command automatically applies Gzip compression and encryption for uploaded files.

[Back to Question 27](#)

Question 28: A. Temporary

Temporary tables last only for the duration of a session and are automatically dropped.

[Back to Question 28](#)

Question 29: D. All of the above

Snowflake Web UI, SQL commands, and integration tools can all be used to create and manage Virtual Warehouses.

[Back to Question 29](#)

Question 30: D. All of the above

When a pipe is recreated, its load history resets, REFRESH defaults to TRUE, and prior files are ignored.

[Back to Question 30](#)

Question 31: D. Business Critical Edition

This edition is the minimum required for customers with regulatory compliance needs.

[Back to Question 31](#)

Question 32: A. Temporary, B. Transient, D. Permanent

Snowflake supports three types of tables: Temporary, Transient, and Permanent. "Provisional" is not valid.

[Back to Question 32](#)

Question 33: A. True

Snowpipe via REST API can only reference external stages as its data source.

[Back to Question 33](#)

Question 34: B. False

Any tool that supports standard JDBC/ODBC can connect to Snowflake using its JDBC/ODBC drivers.

[Back to Question 34](#)

Question 35: A. True

Snowflake allows loading using inline file format options, without requiring a named File Format object.

[Back to Question 35](#)

Question 36: B. False

Tables in Snowflake are independent of the Virtual Warehouse used for loading they can be queried using any active warehouse with permissions.

[Back to Question 36](#)

Question 37: A. VARIANT null values are not the same as SQL NULL values, C. It is recommended to validate the data before loading, D. It is recommended to use staging tables
Frequent single-row DMLs (option B) are discouraged since they are inefficient in Snowflake.

[Back to Question 37](#)

Question 38: A. They are approximately 16MB in size, C. They are immutable
Micro-partitions are fixed in size (around 16MB), compressed automatically, and immutable.

[Back to Question 38](#)

Question 39: A. True

Each query ID is globally unique across all Snowflake deployments and useful for troubleshooting with Support.

[Back to Question 39](#)

Question 40: B. The same query returns precomputed results if data hasnt changed and results accessed in last 24h, D. The 24h timer is renewed each time the exact query executes
Caching ensures faster response as long as results remain valid and accessed within the 24-hour retention.

[Back to Question 40](#)

Question 41: C. Scaling out

Scaling out increases the number of clusters in a multi-cluster warehouse, while scaling up increases the size of a single cluster.

[Back to Question 41](#)

Question 42: A. Snowflake tables are logical representations of underlying physical data
Snowflake manages the physical storage; users work with logical representations.

[Back to Question 42](#)

Question 43: C. Migrate Indexes

Indexes are not used in Snowflake; instead, Snowflake uses micro-partitions for efficient data access.

[Back to Question 43](#)

Question 44: A. Temporary Tables and B. Transient Tables

Both are designed for short-lived or cost-sensitive data storage, reducing costs compared to permanent tables.

[Back to Question 44](#)

Question 45: A. 8KB

Snowflakes maximum compressed row size is 8KB, due to micro-partition storage constraints.

[Back to Question 45](#)

Question 46: A. Databases, C. Warehouses, D. Worksheets

These are the primary sections visible in the Snowflake UI navigation bar.

[Back to Question 46](#)

Question 47: D. VARIANT

The VARIANT data type is specifically designed to store semi-structured data such as JSON, Avro, ORC, or Parquet.

[Back to Question 47](#)

Question 48: B. They roll up and release approximately monthly, but customers can request early release application; C. During a release, new customer requests/queries/connections transparently move over to the newer version

Snowflake upgrades are seamless, transparent, and typically follow a monthly release cycle.

[Back to Question 48](#)

Question 49: A. Quick provisioning of Dev and Test/QA environments; B. Data backups; C. Point in time snapshots

These are the most common scenarios for zero-copy cloning as they require fast and efficient duplication without storage overhead.

[Back to Question 49](#)

Question 50: B. 16

A Medium Warehouse has 8x the capacity of an X-Small, with 16 servers per cluster.

[Back to Question 50](#)

Question 51: B. False

Consumers cannot extend a Data Share to others; only the Data Provider controls the share.

[Back to Question 51](#)

Question 52: B. Is only available to customers with Business Critical Edition; C. Restricts or enables access to specific IP addresses

Network policies are exclusive to Business Critical Edition and allow IP-based access control.

[Back to Question 52](#)

Question 53: B. False

Snowflake does not charge Data Providers for creating shares; costs are incurred only when Data Consumers use the data.

[Back to Question 53](#)

Question 54: D. 24 Hours

Snowflake keeps query results in the result cache for 24 hours, provided the underlying data has not changed.

[Back to Question 54](#)

Question 55: C. SYSADMIN

When a role is dropped, ownership of its objects (like tables) is transferred to the SYSADMIN role by default.

[Back to Question 55](#)

Question 56: A. SnowSQL; B. ODBC

These two connectors are provided directly by Snowflake in the Web UI downloads section.

[Back to Question 56](#)

Question 57: A. UPSERT

Snowflake does not support an explicit UPSERT command. Instead, the MERGE command is used for that functionality.

[Back to Question 57](#)

Question 58: D. At the instance/instant a clone is created, all micro-partitions in the original table and the clone are fully shared.

Zero-copy cloning does not duplicate data; instead, it references the same micro-partitions until changes are made.

[Back to Question 58](#)

Question 59: A. True

By default, the creator of a role becomes its owner and retains ownership until it is explicitly transferred.

[Back to Question 59](#)

Question 60: C. 14 days

Snowflake's query history is retained for 14 days by default in the web interface.

[Back to Question 60](#)

Question 61: D. Set the Minimum Clusters and Maximum Clusters settings to different values

This enables Snowflake to scale out and in automatically, based on workload demand.

[Back to Question 61](#)

Question 62: C. Multi-cluster, shared data

Snowflake's architecture separates compute and storage, enabling multiple compute clusters to access the same shared data simultaneously.

[Back to Question 62](#)

Question 63: C. Auto-resume

D. Auto-suspend

These options help optimize cost and performance by pausing warehouses during inactivity and resuming them automatically when queries are submitted.

[Back to Question 63](#)

Question 64: B. The entire Virtual Warehouse

Auto-suspend and auto-resume apply at the warehouse level, not to individual clusters or queries.

[Back to Question 64](#)

Question 65: A. Temporary

B. Transient

Fail-safe applies only to permanent tables; temporary and transient tables do not have fail-safe.

[Back to Question 65](#)

Question 66: D. Stages

Time Travel applies to tables, schemas, and databases, but not to stages.

[Back to Question 66](#)

Question 67: A. True

Because micro-partition metadata is stored in the cloud services layer, some operations (like metadata-only queries) do not require compute.

[Back to Question 67](#)

Question 68: B. INSERT

D. COPY

Both are non-blocking operations in Snowflake. UPDATE and MERGE can be blocking.

[Back to Question 68](#)

Question 69: C. Snowflake automatically manages the compute required to execute the Pipe's COPY INTO commands

D. Snowpipe keeps track of which files it has loaded

These ensure serverless, continuous ingestion with minimal manual management.

[Back to Question 69](#)

Question 70: C. SECURITYADMIN, ACCOUNTADMIN

These are the critical roles for which MFA is strongly recommended by Snowflake.

[Back to Question 70](#)

Question 71: C. When its provisioning is complete

A Virtual Warehouse can only execute queries after provisioning is fully complete.

[Back to Question 71](#)

Question 72: B. False

Query result sets are private to the user/session that executed them, even if roles overlap.

[Back to Question 72](#)

Question 73: B. False

In a multi-cluster Warehouse, Snowflake automatically distributes queries to clusters. Users do not manually pick clusters.

[Back to Question 73](#)

Question 74: A. True

Pipes can indeed be suspended and resumed using SQL commands.

[Back to Question 74](#)

Question 75: B. Javascript

D. Python

These are the two supported languages for creating UDFs in Snowflake.

[Back to Question 75](#)

Question 76: B. When managing a steady workload

C. When the compute must be available with no delay or lag time

Auto-suspend is best avoided for steady or latency-sensitive workloads.

[Back to Question 76](#)

Question 77: A. Bulk copy from an External Stage

B. Continuous load using Snowpipe REST API

C. The Snowflake Web Interface (UI) data loading wizard

D. Bulk copy from an Internal Stage

(All four are valid Snowflake-supported data loading approaches.)

[Back to Question 77](#)

Question 78: B. The last query using the Warehouse completes

Auto-suspend happens automatically after the warehouse is idle (last query completes).

[Back to Question 78](#)

Question 79: B. False

MFA in Snowflake is supported independently and is not restricted only to SSO scenarios.

[Back to Question 79](#)

Question 80: A. The complexity of each query

B. The `CONCURRENT_QUERY_LIMIT` parameter

Concurrency depends both on query complexity and the systems configured concurrency limit.

[Back to Question 80](#)

Question 81: B. `VALIDATION_MODE=RETURN_ALL_ERRORS` is a parameter of the `COPY` command

D. The `VALIDATION_MODE` option will validate data by the `COPY` statement without completing the load and return errors

Validation mode is a `COPY` option, not for stage creation, and it does not commit data load when validation is active.

[Back to Question 81](#)

Question 82: D. The role must have access to the target schema and the `CREATE TASK` privilege on the schema itself

Task creation requires the proper schema-level privilege and role access, not necessarily global admin role.

[Back to Question 82](#)

Question 83: B. Simplicity

D. Concurrency

E. Performance

These three are the primary expectations customers have from modern cloud data warehouse solutions.

[*Back to Question 83*](#)

Question 84: A. True

Some metadata-only queries can be satisfied from the result/metadata cache without needing a running warehouse.

[*Back to Question 84*](#)

Question 85: A. Concurrency

Scaling out with multi-cluster warehouses is designed to handle more concurrent users and queries, not to improve single-query performance.

[*Back to Question 85*](#)

Question 86: C. Enterprise

Data sharing is supported starting with the Enterprise edition of Snowflake.

[*Back to Question 86*](#)

Question 87: A. True

Each worksheet can be tied to its own role, database, schema, and warehouse context.

[*Back to Question 87*](#)

Question 88: A. True

With external tables and external stages, you can query data in files directly without first loading them into Snowflake tables.

[*Back to Question 88*](#)

Question 89: B. Semi-structured data

The FLATTEN function is specifically used to query semi-structured data types like JSON, VARIANT, and ARRAY.

[*Back to Question 89*](#)