SnowPro Core Practice Questions

Roushan Ara

August 19, 2025

Questions

Question 1

Snowflake provides a mechanism for its customers to override its natural clustering algorithms. This method is:

    A. Micro-partitions

    B. Clustering keys

    C. Key partitions

    D. Clustered partitions

Question 2

Which of the following are valid Snowflake Virtual Warehouse Scaling Policies? (Choose two.)

    A. Custom

    B. Economy

    C. Optimized

    D. Standard

Question 3

True or False: A single database can exist in more than one Snowflake account.

    A. True

    B. False

Question 4

Which of the following roles is recommended to be used to create and manage users and roles?

    A. SYSADMIN

    B. SECURITYADMIN

    C. PUBLIC

    D. ACCOUNTADMIN

Question 5

True or False: Bulk unloading of data from Snowflake supports the use of a SELECT statement.

    A. True

    B. False

Question 6

Select the different types of Internal Stages: (Choose three.)

    A. Named Stage

    B. User Stage

    C. Table Stage

    D. Schema Stage

Question 7

True or False: A customer using SnowSQL / native connectors will be unable to also use the Snowflake Web Interface (UI) unless access to the UI is explicitly granted by support.

    A. True

    B. False

Question 8

Account-level storage usage can be monitored via:

    A. The Snowflake Web Interface (UI) in the Databases section

    B. The Snowflake Web Interface (UI) in the Account → Billing & Usage section

    C. The Information Schema → ACCOUNT_USAGE_HISTORY View

    D. The Account Usage Schema → ACCOUNT_USAGE_METRICS View

Question 9

Credit Consumption by the Compute Layer (Virtual Warehouses) is based on: (Choose two.)

    A. Number of users

    B. Warehouse size

    C. Amount of data processed

    D. Number of clusters for the Warehouse

Question 10

Which statement best describes 'clustering'?

    A. Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions

    B. The database administrator must define the clustering methodology for each Snowflake table

    C. The clustering key must be included on the COPY command when loading data into Snowflake

    D. Clustering can be disabled within a Snowflake account

Question 11

True or False: The COPY command must specify a File Format in order to execute.

    A. True

    B. False

Question 12

Which of the following commands sets the Virtual Warehouse for a session?

    A. COPY WAREHOUSE FROM <config file>

    B. SET WAREHOUSE = <warehouse name>

    C. USE WAREHOUSE <warehouse name>

    D. USE VIRTUAL_WAREHOUSE <warehouse name>

Question 13

Which of the following objects can be cloned? (Choose four.)

    A. Tables

    B. Named File Formats

    C. Schemas

    D. Shares

    E. Databases

    F. Users

Question 14

Which object allows you to limit the number of credits consumed within a Snowflake account?

    A. Account Usage Tracking

    B. Resource Monitor

    C. Warehouse Limit Parameter

    D. Credit Consumption Tracker

Question 15

Snowflake is designed for which type of workloads? (Choose two.)

    A. OLAP (Analytics) workloads

    B. OLTP (Transactional) workloads

    C. Concurrent workloads

    D. On-premise workloads

See answer page 16

Question 16

What are the three layers that make up Snowflake's architecture? (Choose three.)

    A. Compute

    B. Tri-Secret Secure

    C. Storage

    D. Cloud Services

See answer page 16

Question 17

Why would a customer size a Virtual Warehouse from an X-Small to a Medium?

    A. To accommodate more queries

    B. To accommodate more users

    C. To accommodate fluctuations in workload

    D. To accommodate a more complex workload

See answer page 16

Question 18

True or False: Reader Accounts incur no additional Compute costs to the Data Provider since they are simply reading the shared data without making changes.

    A. True

    B. False

See answer page 16

Question 19

Which of the following connectors allow Multi-Factor Authentication (MFA) authorization when connecting? (Choose all that apply.)

A. JDBC

B. SnowSQL

C. Snowflake Web Interface (UI)

D. ODBC

E. Python

Question 20

True or False: Snowflake charges a premium for storing semi-structured data.

A. True

B. False

Question 21

Which of the following statements describes a benefit of Snowflake's separation of compute and storage? (Choose all that apply.)

A. Growth of storage and compute are tightly coupled together

B. Storage expands without the requirement to add more compute

C. Compute can be scaled up or down without the requirement to add more storage

D. Multiple compute clusters can access stored data without contention

Question 22

True or False: It is possible to unload structured data to semi-structured formats such as JSON and Parquet.

A. True

B. False

Question 23

In which layer of its architecture does Snowflake store its metadata statistics?

    A. Storage Layer

    B. Compute Layer

    C. Database Layer

    D. Cloud Services Layer

Question 24

True or False: Data in fail-safe can be deleted by a user or the Snowflake team before it expires.

    A. True

    B. False

Question 25

True or False: Snowflake's data warehouse was built from the ground up for the cloud in lieu of using an existing database or a platform?

    A. True

    B. False

Question 26

Which of the following statements are true of Virtual Warehouses? (Choose all that apply.)

    A. Customers can change the size of the Warehouse after creation

    B. A Warehouse can be resized while running

    C. A Warehouse can be configured to suspend after a period of inactivity

    D. A Warehouse can be configured to auto-resume when new queries are submitted

Question 27

The PUT command: (Choose two.)

    A. Automatically creates a File Format object

    B. Automatically uses the last Stage created

    C. Automatically compresses files using Gzip

    D. Automatically encrypts files

Question 28

Which type of table corresponds to a single Snowflake session?

    A. Temporary

    B. Transient

    C. Provisional

    D. Permanent

Question 29

Which interfaces can be used to create and/or manage Virtual Warehouses?

    A. The Snowflake Web Interface (UI)

    B. SQL commands

    C. Data integration tools

    D. All of the above

Question 30

When a Pipe is recreated using the CREATE OR REPLACE PIPE command:

    A. The Pipe load history is reset to empty

    B. The REFRESH parameter is set to TRUE

    C. Previously loaded files will be ignored

    D. All of the above

Question 31

What is the minimum Snowflake edition that customers planning on storing protected information in Snowflake should consider for regulatory compliance?

A. Standard

B. Premier

C. Enterprise

D. Business Critical Edition

Question 32

Select the three types of tables that exist within Snowflake. (Choose three.)

A. Temporary

B. Transient

C. Provisional

D. Permanent

Question 33

True or False: Snowpipe via REST API can only reference External Stages as source.

A. True

B. False

Question 34

True or False: A third-party tool that supports standard JDBC or ODBC but has no Snowflake-specific driver will be unable to connect to Snowflake.

A. True

B. False

Question 35

True or False: It is possible to load data into Snowflake without creating a named File Format object.

A. True

B. False

Question 36

True or False: A table in Snowflake can only be queried using the Virtual Warehouse that was used to load the data.

    A. True

    B. False

Question 37

Which of the following statements are true of Snowflake data loading? (Choose three.)

    A. VARIANT null values are not the same as SQL NULL values

    B. It is recommended to do frequent, single row DMLs

    C. It is recommended to validate the data before loading into the Snowflake target table

    D. It is recommended to use staging tables to manage MERGE statements

Question 38

Which statements are true of micro-partitions? (Choose two.)

    A. They are approximately 16MB in size

    B. They are stored compressed only if COMPRESS=TRUE on Table

    C. They are immutable

    D. They are only encrypted in the Enterprise edition and above

Question 39

True or False: Query IDs are unique across all Snowflake deployments and can be used in communication with Snowflake Support to help troubleshoot issues.

    A. True

    B. False

Question 40

A deterministic query is run at 8am, takes 5 minutes, and the results are cached. Which of the following statements are true? (Choose two.)

   A. The exact query will ALWAYS return the precomputed result set for the RESULT_CACHE_ACTIVE = time period

   B. The same exact query will return the precomputed results if the underlying data hasnt changed and the results were last accessed within previous 24 hour period

   C. The same exact query will return the precomputed results even if the underlying data has changed as long as the results were last accessed within the previous 24 hour period

   D. The 24 hour timer on the precomputed results gets renewed every time the exact query is executed

Question 41

Increasing the maximum number of clusters in a Multi-Cluster Warehouse is an example of:

   A. Scaling rhythmically

   B. Scaling max

   C. Scaling out

   D. Scaling up

Question 42

Which statement best describes Snowflake tables?

   A. Snowflake tables are logical representations of underlying physical data

   B. Snowflake tables are the physical instantiation of data loaded into Snowflake

   C. Snowflake tables require that clustering keys be defined to perform optimally

   D. Snowflake tables are owned by a user

Question 43

Which item in the Data Warehouse migration process does not apply in Snowflake?

   A. Migrate Users

   B. Migrate Schemas

   C. Migrate Indexes

   D. Build the Data Pipeline

Question 44

Snowflake provides two mechanisms to reduce data storage costs for short-lived tables. These mechanisms are: (Choose two.)

A. Temporary Tables

B. Transient Tables

C. Provisional Tables

D. Permanent Tables

Question 45

What is the maximum compressed row size in Snowflake?

A. 8KB

B. 16MB

C. 50MB

D. 4000GB

Question 46

Which of the following are main sections of the top navigation of the Snowflake Web Interface (UI)? (Choose three.)

A. Databases

B. Tables

C. Warehouses

D. Worksheets

Question 47

What is the recommended Snowflake data type to store semi-structured data like JSON?

A. VARCHAR

B. RAW

C. LOB

D. VARIANT

Question 48

Which of the following statements are true of Snowflake releases: (Choose two.)

A. They happen approximately weekly

B. They roll up and release approximately monthly, but customers can request early release application

C. During a release, new customer requests/queries/connections transparently move over to the newer version

D. A customer is assigned a 30 minute window (that can be moved anytime within a week) during which the system will be unavailable and customer is upgraded

Question 49

Which of the following are common use cases for zero-copy cloning? (Choose three.)

A. Quick provisioning of Dev and Test/QA environments

B. Data backups

C. Point in time snapshots

D. Performance optimization

Question 50

If a Small Warehouse is made up of 2 servers/cluster, how many servers/cluster make up a Medium Warehouse?

A. 4

B. 16

C. 32

D. 128

Question 51

True or False: When a data share is established between a Data Provider and a Data Consumer, the Data Consumer can extend that data share to other Data consumers.

A. True

B. False

Question 52

Which is true of Snowflake network policies? A Snowflake network policy: (Choose two.)

   A. Is available to all Snowflake Editions

   B. Is only available to customers with Business Critical Edition

   C. Restricts or enables access to specific IP addresses

   D. Is activated using an ALTER DATABASE command

Question 53

True or False: Snowflake charges additional fees to Data Providers for each Share they create.

   A. True

   B. False

Answers

Question 1: B. Clustering keys
They allow customers to explicitly define clustering for performance optimization.
Back to Question 1

Question 2: B. Economy **and** D. Standard
These are the only valid Snowflake Virtual Warehouse scaling policies.
Back to Question 2

Question 3: B. False
A Snowflake database belongs to a single account and cannot exist in more than one account.
Back to Question 3

Question 4: B. SECURITYADMIN
This role is specifically responsible for creating and managing users and roles.
Back to Question 4

Question 5: A. True
Bulk unloading supports the use of a SELECT statement to filter and transform data before writing to files.
Back to Question 5

Question 6: A. Named Stage, B. User Stage, and C. Table Stage
These are the three types of internal stages in Snowflake.
Back to Question 6

Question 7: B. False
SnowSQL and native connectors can be used independently of the Web UI; UI access does not require explicit support enablement.
Back to Question 7

Question 8: D. The Account Usage Schema → ACCOUNT_USAGE_METRICS View
This view provides account-level storage usage details.
Back to Question 8

Question 9: B. Warehouse size **and** D. Number of clusters for the Warehouse
Credit consumption depends on warehouse size and the number of clusters in multi-cluster warehouses.
Back to Question 9

Question 10: A. Clustering represents the way data is grouped together and stored within Snowflake's micro-partitions
Clustering is Snowflake's method of organizing micro-partition storage.
Back to Question 10

Question 11: A. True
The COPY command requires a file format (or defaults) to execute properly.
Back to Question 11

Question 12: C. USE WAREHOUSE <warehouse name>
This command sets the active virtual warehouse for the session.

Question 13: A. Tables, C. Schemas, D. Shares, E. Databases
These Snowflake objects can be cloned. Named file formats and users cannot.

Question 14: B. Resource Monitor
Resource monitors are used to limit and control credit consumption within a Snowflake account.

Question 15: A. OLAP (Analytics) workloads **and** C. Concurrent workloads
Snowflake is designed for analytics and concurrent workloads, not OLTP or on-premise.

Question 16: A. Compute, C. Storage, D. Cloud Services
These are the three layers of Snowflakes architecture.

Question 17: B. To accommodate more users
Larger virtual warehouses support more users running queries concurrently.

Question 18: B. False
Reader Accounts still generate compute costs for the provider since queries executed on shared data consume compute resources.

Question 19: C. Snowflake Web Interface (UI) and B. SnowSQL
MFA is supported through the web interface and SnowSQL client, but not directly via JDBC/ODBC/Python.

Question 20: B. False
Snowflake does not charge a premium for semi-structured data; it is stored in VARIANT columns without extra fees beyond normal storage.

Question 21: B. Storage expands without adding compute, C. Compute scales independently of storage, D. Multiple clusters can access data without contention
These are the key benefits of separating compute from storage in Snowflake.

Question 22: A. True
Snowflake supports unloading structured data into semi-structured formats such as JSON and Parquet.

Question 23: D. Cloud Services Layer
Snowflake stores and manages metadata statistics in the Cloud Services layer.

Question 24: B. False
Fail-safe is managed entirely by Snowflake and cannot be modified or deleted by users or Snowflake staff before expiration.

Question 25: A. True
Snowflake was designed from the ground up as a cloud-native data warehouse platform, not adapted from an on-premise system.

Question 26: A. Customers can change size after creation, C. Suspend after inactivity, D. Auto-resume on query submission
Warehouses cannot be resized dynamically while running (option B is false).

Question 27: C. Automatically compresses files using Gzip, D. Automatically encrypts files
The PUT command automatically applies Gzip compression and encryption for uploaded files.

Question 28: A. Temporary
Temporary tables last only for the duration of a session and are automatically dropped.
Question 29: D. All of the above
Snowflake Web UI, SQL commands, and integration tools can all be used to create and manage Virtual Warehouses.
Question 30: D. All of the above
When a pipe is recreated, its load history resets, REFRESH defaults to TRUE, and prior files are ignored.
Question 31: D. Business Critical Edition
This edition is the minimum required for customers with regulatory compliance needs.
Question 32: A. Temporary, B. Transient, D. Permanent
Snowflake supports three types of tables: Temporary, Transient, and Permanent. "Provisional" is not valid.
Question 33: A. True
Snowpipe via REST API can only reference external stages as its data source.
Question 34: B. False
Any tool that supports standard JDBC/ODBC can connect to Snowflake using its JDBC/ODBC drivers.
Question 35: A. True
Snowflake allows loading using inline file format options, without requiring a named File Format object.

Question 36: B. False
Tables in Snowflake are independent of the Virtual Warehouse used for loading they can be queried using any active warehouse with permissions.
Question 37: A. VARIANT null values are not the same as SQL NULL values, C. It is recommended to validate the data before loading, D. It is recommended to use staging tables Frequent single-row DMLs (option B) are discouraged since they are inefficient in Snowflake.
Question 38: A. They are approximately 16MB in size, C. They are immutable
Micro-partitions are fixed in size (around 16MB), compressed automatically, and immutable.
Question 39: A. True
Each query ID is globally unique across all Snowflake deployments and useful for troubleshooting with Support.
Question 40: B. The same query returns precomputed results if data hasnt changed and results accessed in last 24h, D. The 24h timer is renewed each time the exact query executes Caching ensures faster response as long as results remain valid and accessed within the 24-hour retention.
Question 41: C. Scaling out
Scaling out increases the number of clusters in a multi-cluster warehouse, while scaling up increases the size of a single cluster. **Question 42:** A. Snowflake tables are logical representations of underlying physical data
Snowflake manages the physical storage; users work with logical representations.
Question 43: C. Migrate Indexes
Indexes are not used in Snowflake; instead, Snowflake uses micro-partitions for efficient data access.
Question 44: A. Temporary Tables and B. Transient Tables
Both are designed for short-lived or cost-sensitive data storage, reducing costs compared to permanent tables.
Question 45: A. 8KB
Snowflakes maximum compressed row size is 8KB, due to micro-partition storage constraints.
Question 46: A. Databases, C. Warehouses, D. Worksheets
These are the primary sections visible in the Snowflake UI navigation bar.
Question 47: D. VARIANT
The VARIANT data type is specifically designed to store semi-structured data such as JSON, Avro, ORC, or Parquet.
Question 48: B. They roll up and release approximately monthly, but customers can request early release application; C. During a release, new customer requests/queries/connections transparently move over to the newer version
Snowflake upgrades are seamless, transparent, and typically follow a monthly release cycle.
Question 49: A. Quick provisioning of Dev and Test/QA environments; B. Data backups; C. Point in time snapshots
These are the most common scenarios for zero-copy cloning as they require fast and efficient duplication without storage overhead.
Question 50: B. 16
A Medium Warehouse has 8x the capacity of an X-Small, with 16 servers per cluster.
Question 51: B. False
Consumers cannot extend a Data Share to others; only the Data Provider controls the share.

Question 52: B. Is only available to customers with Business Critical Edition; C. Restricts or enables access to specific IP addresses
Network policies are exclusive to Business Critical Edition and allow IP-based access control.
Question 53: B. False
Snowflake does not charge Data Providers for creating shares; costs are incurred only when Data Consumers use the data.